

# Medium

 Search[Write](#)[Sign up](#)[Sign in](#)

♦ Member-only story

# Mastering the Deep Learning Interview: Top 35 Questions and Expert Answers

Ritesh Gupta · [Follow](#)

Published in Artificial Intelligence in Plain English · 13 min read · Jan 23, 2023

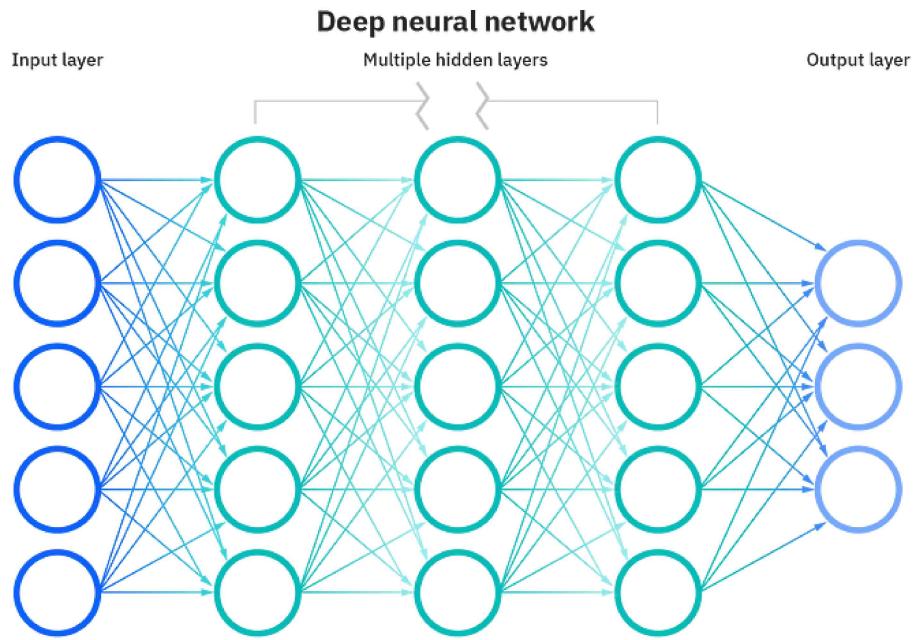
745

10



## What is Deep Learning?

Deep learning is a subset of machine learning in which artificial neural networks, algorithms inspired by the human brain, learn from large amounts of data. These models can be used for a variety of tasks such as image and speech recognition, natural language processing, and decision-making. Deep learning has led to breakthroughs in many industries, from self-driving cars to medical diagnosis. It has become one of the most exciting and rapidly growing areas of artificial intelligence research and development.



The mathematical equation for a single artificial neuron in a neural network is given by the following:

$$y = f(w_1x_1 + w_2x_2 + \dots + w_nx_n + b)$$

**where:**

$y$  is the output of the neuron

$x_1, x_2, \dots, x_n$  are the inputs to the neuron

$w_1, w_2, \dots, w_n$  are the weights associated with each input

$b$  is the biased term

$f$  is the activation function, which is a non-linear function that determines the output of the neuron

For a neural network with multiple layers, this equation is applied for each neuron in each layer, and the output of one layer becomes the input of the next.

In addition to this, there are many types of activation functions such as sigmoid, tanh, relu, and many more. They can be chosen based on the specific problem and architecture.

The training process of the neural network is an optimization problem. The goal is to find the weights and biases that minimize the error between the predicted output and the actual output. This is typically done using an optimization algorithm such as stochastic gradient descent (SGD) or backpropagation.

1: What is deep learning and how does it differ from traditional machine learning?

A: Deep learning is a subset of machine learning that involves training artificial neural networks to perform tasks that are typically done by humans, such as image and speech recognition, natural language processing, and decision-making. It differs from traditional machine learning in that it uses multiple layers of artificial neurons, known as a deep neural network, to extract and process information from large and complex data sets.

2. How does weight initialization impact the performance of a deep learning model?

A: Weight initialization is an important step in training a deep learning model, as it determines the starting point of the optimization process. If the weights are initialized with small random values, the model will have a hard time learning from the data, as the gradients will be too small. On the other hand, if the weights are initialized with large random values, the model will converge quickly, but it may not find the global minimum of the loss function. Choosing an appropriate weight initialization method, such as He initialization or Glorot initialization, can help improve the performance of the model.

3. Can you explain the difference between a perceptron and a sigmoid neuron?

A: A perceptron is a single-layer neural network that uses a linear activation function and is used for binary classification. A sigmoid neuron is a single-layer neural network that uses a sigmoid activation function and can be used for both binary and multi-class classification.

4. Can you explain how backpropagation works in a neural network?

A: Backpropagation is the process by which an artificial neural network learns from its errors. It is used to adjust the weights of the network's

neurons in order to minimize the error between the network's predictions and the actual output. The process starts by forwarding the input through the network to calculate the output. Then the error is calculated by comparing the output to the desired output. The error is then propagated back through the network and the weights are adjusted in such a way that the error is minimized.

#### 5. What are some common activation functions used in neural networks?

A: Some common activation functions used in neural networks include sigmoid, ReLU (Rectified Linear Unit), and tanh (hyperbolic tangent). Sigmoid maps any input value to a value between 0 and 1, and is often used in the output layer of a binary classification problem. ReLU maps any input value less than 0 to 0 and any value greater than or equal to 0 to the input value itself, which helps to alleviate the vanishing gradient problem. tanh maps any input value to a value between -1 and 1, and is often used in the hidden layers of a network.

#### 6. How do you prevent overfitting in a neural network?

A: Overfitting occurs when a model is trained too well on the training data and performs poorly on new, unseen data. There are several techniques that can be used to prevent overfitting in a neural network, such as regularization, early stopping, dropout, and cross-validation. Regularization involves adding a penalty term to the loss function to discourage the network from assigning large weights to certain neurons. Early stopping involves monitoring the performance of the network on a validation set and stopping the training when performance starts to plateau. Dropout involves randomly turning off a certain percentage of neurons during training, in order to reduce the dependency of the network on any particular neuron. Cross-validation involves dividing the data into multiple subsets and training and evaluating the model on different subsets, in order to get a better estimate of the model's performance on unseen data.

#### 7. How do you decide the number of layers and neurons in a neural network?

A: The number of layers and neurons in a neural network is known as the architecture of the network. The architecture of a network is typically determined through a process of trial and error. A common approach is to start with a small number of layers and neurons and gradually increase them until the performance of the network on a validation set stops improving. Another approach is to use techniques such as grid search and

random search to systematically explore different architectures and find the one that performs best.

8. How do you handle missing data in a neural network?

A: Handling missing data in a neural network can be a challenge. One common approach is to simply remove any samples that contain missing data, but this may lead to a loss of information. Another approach is to impute the missing values using techniques such as mean imputation, median imputation, or multiple imputation. A more sophisticated approach is to use a neural network to predict the missing values based on the other available information.

9: Can you explain the concept of transfer learning in deep learning?

A: Transfer learning is the process of using a pre-trained neural network as a starting point for a new task. It involves taking the weights of a network that has already been trained on a large dataset and using them as the initial weights for a new network that is being trained on a smaller dataset. This can be useful when there is a shortage of data for a particular task, or when the new task is similar to the original task. By using the pre-trained network as a starting point, the new network can learn from the features that were learned by the original network, and thus requires less data and less training time to achieve good performance.

10. How do you evaluate the performance of a deep learning model?

A: Evaluating the performance of a deep learning model can be challenging, as it depends on the specific task and the characteristics of the data. Common metrics used to evaluate the performance of a deep learning model include accuracy, precision, recall, F1 score, and ROC-AUC. Another common method is to evaluate the model on unseen test data and compare the performance with other models. It's also important to consider the business impact of the model and how well it meets the needs of the end-users.

11. Can you explain the difference between a convolutional neural network and a recurrent neural network?

A: A convolutional neural network (CNN) is a type of neural network that is used for image and video processing. It uses convolutional layers to learn features of the data and reduce dimensionality. A recurrent neural network

(RNN) is a type of neural network that is used for sequential data, such as time series or natural language. It uses recurrent connections to allow information to flow through the network and maintain a memory of previous inputs.

12. What is the difference between overfitting and underfitting?

A: Overfitting occurs when a model is trained too well on the training data, and as a result, it performs poorly on unseen data. This is because the model has learned the noise in the training data. Underfitting occurs when a model is not trained well enough on the training data and as a result, it also performs poorly on unseen data. This is usually because the model is too simple and cannot learn the underlying patterns in the data.

13. What is the purpose of Dropout in a neural network?

A: Dropout is a regularization technique used to prevent overfitting in neural networks by randomly dropping out neurons during training. This helps to reduce the co-adaptation of neurons, which can occur when neurons in a network become too reliant on one another.

14. How does batch normalization work?

A: Batch normalization is a technique used to normalize the activations of a neural network, which helps to stabilize the training process and speed up convergence. It works by normalizing the activations of each layer by subtracting the batch mean and dividing by the batch standard deviation. This helps to ensure that the inputs to each layer have a similar scale, which can help to prevent the vanishing or exploding gradients problem.

15. What is the role of the activation function in a neural network?

A: The activation function in a neural network is used to introduce non-linearity into the network. Without non-linearity, a neural network would just be a linear model and would not be able to learn complex relationships in the data. Activation functions, such as the sigmoid, ReLU, and tanh, are used to introduce non-linearity into the network by squashing the input to a certain range, such as between 0 and 1 or between -1 and 1.

16. What is the difference between a deep neural network and a shallow neural network?

A: A deep neural network is a neural network with many layers, typically more than 3 layers. A shallow neural network, on the other hand, is a neural network with only one or two layers. In general, deep neural networks are able to learn more complex features and representations of the data than shallow neural networks, but they are also more computationally expensive and can be more difficult to train.

17. What is the difference between gradient descent and stochastic gradient descent?

A: Gradient descent is an optimization algorithm used to find the minimum of a function by moving in the opposite direction of the gradient of the function. Stochastic gradient descent (SGD) is an optimization algorithm that uses random samples of the data to estimate the gradient of the function, rather than using the full dataset. SGD is an approximation of gradient descent and is often used when the dataset is too large to be processed in memory.

18. What are recurrent neural networks?

Recurrent neural networks (RNNs) are a type of neural network that are particularly well-suited for sequential data, such as time series data or natural language text. RNNs use a feedback loop, or recurrence, to process sequences of variable length and maintain an internal memory of past inputs.

19. What is the difference between a validation set and a test set in deep learning?

A validation set is used to tune the hyperparameters of a model and evaluate its performance during the training process. A test set is used to evaluate the final performance of the model, after the training process is completed and the hyperparameters have been chosen.

20. What is the difference between a feedforward neural network and a recurrent neural network?

A feedforward neural network is a type of neural network where the data flows in one direction, from input layer to output layer, without loops. In contrast, a recurrent neural network has a feedback loop, or recurrence, which allows it to maintain an internal memory of past inputs and process sequences of variable length.

21. Can you explain the concept of weight initialization in deep learning?

Weight initialization is the process of setting the initial values of the weights in a neural network. Proper weight initialization is important because it can affect the speed and stability of the training process. There are several methods for weight initialization, such as random initialization and Xavier initialization, which are designed to prevent the vanishing or exploding gradients problem.

22. What is the difference between a softmax function and a sigmoid function?

A softmax function is a generalization of the sigmoid function, commonly used in the output layer of a neural network for multi-class classification problem. The softmax function produces a probability distribution over the classes, while the sigmoid function produces a probability value between 0 and 1 for a binary classification problem.

23. Can you explain the concept of ensemble learning in deep learning?

Ensemble learning is a technique that combines the predictions of multiple models to improve overall performance. This can be achieved by training multiple models independently and then averaging their predictions, or by training a single model that combines the predictions of multiple sub-models. Ensemble learning can increase the robustness and generalization of the model.

24. Can you explain the concept of weight decay in deep learning?

Weight decay is a regularization technique that aims to prevent overfitting by adding a penalty term to the loss function, based on the magnitude of the weights. It works by encouraging the weights to be small, which in turn reduces the complexity of the model. This can be achieved by adding a term to the loss function that is proportional to the sum of squares of the weights.

25. What is the difference between L1 and L2 regularization in deep learning?

L1 regularization is a technique that adds a penalty term to the loss function that is proportional to the absolute value of the weights, while L2 regularization is a technique that adds a penalty term to the loss function that is proportional to the square of the weights. L1 regularization tends to

drive the weights of the model towards zero, resulting in a sparse model with fewer non-zero weights, while L2 regularization tends to drive the weights of the model towards smaller values but not necessarily zero.

26. What is the role of learning rate in deep learning?

The learning rate is a hyperparameter that determines the step size at which the optimizer updates the weights of the model during training. A larger learning rate can lead to faster convergence, but also a higher risk of overshooting the optimal solution. A smaller learning rate can lead to slower convergence, but also a lower risk of overshooting the optimal solution.

27. Can you explain the concept of early stopping in deep learning?

Early stopping is a technique used to prevent overfitting by stopping the training process before the model reaches the point of overfitting. It works by monitoring the performance of the model on a validation set during training, and stopping the training when the performance on the validation set starts to deteriorate.

28. Can you explain the concept of data augmentation in deep learning?

Data augmentation is a technique used to artificially increase the size of a dataset by applying a series of operations to the existing data. These operations can include random cropping, flipping, rotation, scaling, and other forms of image processing. Data augmentation can help to reduce overfitting and improve the generalization of a model by creating new, slightly different versions of the original data.

29. What is the difference between batch normalization and layer normalization?

Batch normalization normalizes the activations of the previous layer by the mean and standard deviation of the mini-batch. Layer normalization normalizes the activations of the previous layer by the mean and standard deviation of the entire layer. While both techniques are used to stabilize the training process and improve the performance of a model, batch normalization is typically used in feedforward neural networks while layer normalization is used in recurrent neural networks.

30. Can you explain the concept of the long short-term memory (LSTM) network?

The long short-term memory (LSTM) network is a type of recurrent neural network that is particularly well-suited for sequential data, such as time series data or natural language text. LSTMs use a special type of memory cell, called an LSTM cell, which can selectively store and retrieve information from the past. This allows LSTMs to maintain an internal memory of past inputs and process sequences of variable length.

31. What is the difference between a generative model and a discriminative model in deep learning?

A generative model is a type of model that learns to generate new data that is similar to the training data, while a discriminative model is a type of model that learns to distinguish between different classes of data. Generative models are typically used to generate new data, such as images or text, while discriminative models are typically used for classification tasks.

32. Can you explain the concept of the Generative Adversarial Networks (GANs)?

A Generative Adversarial Networks (GANs) is a type of model that consists of two neural networks, a generator and a discriminator. The generator network is trained to generate new data that is similar to the training data, while the discriminator network is trained to distinguish between the generated data and the real data. The two networks are trained together, with the generator trying to produce data that can fool the discriminator and the discriminator trying to correctly identify the generated data. GANs are widely used for tasks such as image synthesis, text generation, and other forms of data generation.

33. What is the difference between a deep belief network and a deep neural network?

A deep belief network (DBN) is a type of generative model that consists of multiple layers of restricted Boltzmann machines (RBMs). DBNs are typically trained in a greedy layer-wise manner, where each layer is trained to reconstruct the activations of the previous layer. Deep neural networks (DNNs) are feedforward neural networks with multiple layers that can be trained using techniques such as backpropagation. While both DBNs and DNNs are deep architectures, DBNs are generative models and DNNs are discriminative models.

34. Can you explain the concept of reinforcement learning in deep learning?

Reinforcement learning is a type of learning where an agent learns to make decisions in an environment by performing actions and receiving rewards or penalties. The agent's goal is to learn a policy that maximizes the expected cumulative reward over time. This can be achieved by using techniques such as Q-learning and policy gradients, which involve training a neural network to approximate the value function or the policy.

35. Can you explain the concept of the Attention Mechanism in deep learning?

Attention mechanism is a technique used to focus on specific parts of the input in a neural network. It allows the model to selectively attend to different parts of the input, such as different words in a sentence or different regions in an image, and to weight the importance of each part differently. Attention mechanisms are widely used in tasks such as machine translation, image captioning, and other tasks that involve sequential data or multi-modal data.

## Thanks for Reading!

*If you enjoyed this, [follow](#) me to never miss another article on data science guides, tricks and tips, life lessons, and more!*

### In Plain English

*Thank you for being a part of the [In Plain English](#) community! Before you go:*

- Be sure to clap and follow the writer 
- Follow us: [X](#) | [LinkedIn](#) | [YouTube](#) | [Discord](#) | [Newsletter](#)
- Visit our other platforms: [CoFeed](#) | [Differ](#)
- More content at [PlainEnglish.io](#)

Deep Learning

Interview Questions

Machine Learning

Data Science

TensorFlow



745



10





## Written by Ritesh Gupta

3.4K Followers · Writer for Artificial Intelligence in Plain English

[Follow](#)

Data Scientist, I write Article on Machine Learning| Deep Learning| NLP | Open CV |  
AI Lover ❤️

### More from Ritesh Gupta and Artificial Intelligence in Plain English



Ritesh Gupta

## Can You Handle These 25 Toughest Data Science Interview Questions?

The role of a Data Scientist demands a unique blend of skills, including statistics, machine...

★ Sep 25 46



**2 Query Execution**

movie_title	revenue	review	genre
Shang-Chi	432.2	"solid film..."	Action
Titanic	2257.8	"still best..."	Romance
Titanic	2257.8	"a guilty..."	Romance
...	...	...	...

**3 Answer Generation**

movie_title	revenue	review	genre
Titanic	2257.8	"still best..."	Romance
Titanic	2257.8	"a guilty..."	Romance
...	...	...	...

Pavan Ema... in Artificial Intelligence in Plain Engl...

## Goodbye, Text2SQL: Why Table-Augmented Generation (TAG) is...

Exploring the Future of Natural Language Queries with Table-Augmented Generation.

★ Sep 11 1.1K 17



Andrew Be... in Artificial Intelligence in Plain Engl...

## New KILLER ChatGPT Prompt—The “Playoff Method”

Super powerful prompt for ChatGPT—01  
Preview

★ Sep 27 1.8K 43



Ritesh Gupta

## 14 Life Changing Lessons From Chanakya Niti everyone should...

Who is Chanakya?

★ Jan 30, 2023 110 1



[See all from Ritesh Gupta](#)

[See all from Artificial Intelligence in Plain English](#)

## Recommended from Medium



Axel Casas, PhD Candidate in Stackademic

### This CV Got Me A Data Analyst Job At Disney Without Experience

And without a Bootcamp or Computer Science degree

May 4 · 41 saves · 1 comment



$$P_N = \frac{\alpha}{360^\circ} \cdot \pi \cdot r^2$$
$$\frac{P_N}{\pi \cdot r^2} = \frac{\alpha}{360^\circ}$$
$$\frac{t}{2 \cdot \pi \cdot r} = \frac{\alpha}{360^\circ}$$

Egor Howell in Towards Data Science

### How to Learn the Math Needed for Data Science

A breakdown of the three fundamental math fields required for data science: statistics,...

Mar 4 · 2.2K saves · 13 comments



## Lists



### Predictive Modeling w/ Python

20 stories · 1582 saves



### Practical Guides to Machine Learning

10 stories · 1924 saves



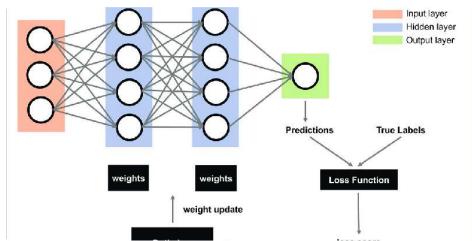
### Natural Language Processing

1747 stories · 1336 saves



### data science and AI

40 stories · 262 saves

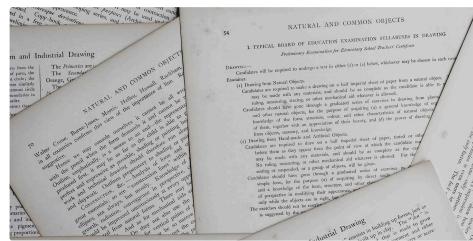


Abhijith Neil Abraham

## A very basic roadmap for starting with Machine Learning and...

Table of Contents:

Apr 24 73 2

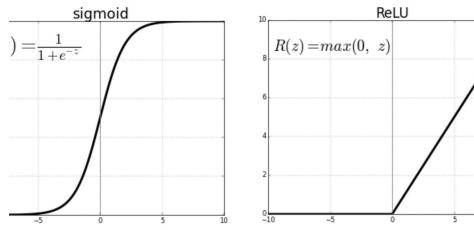


Youssef Hosni in Towards AI

## How to Read Machine Learning Papers Effectively

The field of machine and deep learning is evolving very fast, and there are new research...

Oct 9, 2022 1.8K 12



Shivansh Srivastava

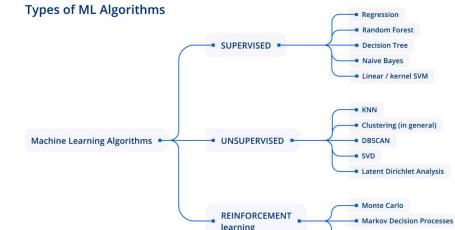
## Understanding the Difference Between ReLU and Sigmoid...

In the realm of deep learning, activation functions play a crucial role in shaping the...

Apr 18 1



### Types of ML Algorithms



John Vastola

## 10 Must-Know Machine Learning Algorithms for Data Scientists

Machine learning is the science of getting computers to act without being explicitly...

Dec 6, 2022 473 7



[See more recommendations](#)

